

## Part 3 – Database Design

### Project Title: Smart Radar Traffic Monitoring System

---

#### Overview

The database design for the *Smart Radar Traffic Monitoring System* is built to efficiently handle large-scale traffic data generated by radar sensors.

It supports **real-time streaming** and **batch processing** workflows, enabling quick analysis, alerting, and visualization.

The system uses a **two-layer storage architecture**:

1. **Azure Data Lake** — for storing raw and processed Parquet files.
  2. **PostgreSQL Data Warehouse** — for structured queries and Power BI dashboards.
- 

#### Database Objectives

- Store, clean, and organize incoming radar data.
  - Log all types of traffic violations and real-time alerts.
  - Provide fast access for analytics and dashboard queries.
  - Ensure data integrity and scalability for large volumes.
- 

#### Main Tables

Table Name	Description	Key Fields
<b>vehicles</b>	Stores all raw vehicle detection data collected by simulated radars.	plate_number, color, speed, location, timestamp, seat_belt, phone_usage
<b>violations</b>	Records detected violations, such as speeding or phone usage.	violation_id, plate_number, violation_type, timestamp, location


Table Name	Description	Key Fields
alerts	Stores alerts triggered in real-time by the streaming system (e.g., wanted vehicles).	alert_id, plate_number, alert_message, severity, alert_time
aggregated_reports	Stores daily or weekly summarized data for dashboard use.	report_date, total_vehicles, total_violations, avg_speed, top_location

### Relationships (ERD Summary)

- One Vehicle → Many Violations**  
Each vehicle may appear multiple times in the system if it commits more than one violation.
- One Violation → May Trigger One Alert**  
A severe violation (e.g., a wanted vehicle) can generate an alert notification.
- Aggregated Reports** summarize historical data from **vehicles** and **violations** tables.

---

### Example ERD Structure




---

### Data Flow Integration

**1. Data Simulation & Ingestion:**

Python scripts generate synthetic radar data → sent to Kafka topics.

**2. Batch Processing:**

PySpark reads Kafka data → cleans & writes to Azure Data Lake (Parquet).

**3. Streaming Processing:**

Real-time Spark job consumes Kafka stream → detects violations → writes alerts to PostgreSQL.

**4. Visualization Layer:**

Power BI connects to PostgreSQL → generates visual reports and insights.

---

**Data Integrity & Quality Rules**

- Records with missing or invalid plate\_number or negative speed are removed.
- Duplicated events are filtered using unique event IDs.
- All timestamps are standardized to UTC format.
- Referential integrity between **violations** and **vehicles** is enforced.